
The Instrument Moves: Unrecorded Serving Changes in Agent Benchmarks

Jingyuan Yi

Independent Researcher

jingyuay@alumni.cmu.edu

Abstract

Agent benchmarks are instruments for measuring capability, and their outputs are read as measurements. We pinned every control a mobile-agent benchmark exposes (model identifier, serving provider, quantization, temperature, seed, image digest, task-set commit, one fresh container per repetition, serial execution) and ran ten repetitions. Three things followed. First, `temperature=0` with a fixed seed did not produce determinism on either endpoint we tried. Our first check said otherwise: a short prompt came back byte-identical twice, on an endpoint that turned out not to be deterministic. Second, ten and a half hours in, the pinned endpoint’s behaviour changed. Ten tasks got worse and none got better. Median completion tokens per action fell by a fifth at a single round boundary and stayed there. We observed the endpoint’s behaviour, not its configuration: the shift was sudden, specific to that endpoint, absent from a control, and invisible to every field a leaderboard records. Third, for this agent, this endpoint and this reweighted task sample, two single-run scores need to differ by roughly eight points before the gap exceeds what repetition alone produces. The token signal flagged the shift at the round level; the score did not. The signal costs one integer per request to log. We argue that an evaluation harness should be treated as a verification system rather than a score generator. Data, analysis code and figure scripts: <https://github.com/Universeyi/instrument-moves>.

1 Introduction

Leaderboards for agent benchmarks are read as measurements. The gap between adjacent entries is cited as evidence that one system is better than another, and differences of two or three points routinely carry that weight. On the board we study, 37 of 40 entries are the result of a single run, and no entry reports an interval.

A benchmark that drives a real operating system through a real user interface against real network services will have run-to-run variance. That is the price of not being a static test set, and it is a price worth paying: the alternative measures something much less like deployment. What has been missing is a number for it, and a way to tell when the number has changed.

We pinned every control the tooling of MobileWorld [Kong et al., 2026] exposes and ran ten repetitions over 29 GUI-only tasks, 290 runs in total. We set out to measure a noise floor. We found something larger: the configuration was not, in the event, fixed, and nothing in the tooling or the reporting format would have told us.

The claim. A benchmark is an instrument for measuring agent capability. This instrument moves while nobody is recording, and current reporting practice contains no mechanism that would reveal it.

Table 1: The pinned configuration. Nothing under the benchmark was modified; the custom agent is supplied through the harness’s existing `-agent-type` hook, so reproducing this requires no fork.

model	moonshotai/kimi-k2.5 via an inference router
serving endpoint	provider pinned, quantization int4, allow_fallbacks: false
sampling	temperature=0, seed=42, max_tokens=2048 (the agent default)
container image	tag v1.4; the full digest is recorded on every row and given in Appendix E
task-set commit	83e7b8fc
execution	-max-concurrency 1, no task shuffling, one fresh container per repetition round
round budget	-max-round 50, -auto-retry 10
tasks	29 GUI-only, stratified from the published trajectory bundles
pass threshold	score > 0.99, the benchmark’s own

Contributions.

- **A methodological error in determinism checks.** Identical requests at temperature=0 with a fixed seed returned different completions on both endpoints we tested, though both advertise seed support. Our own check used a short, tightly constrained prompt and returned byte-identical output twice, which looked like proof of determinism. Longer requests to the same endpoint came back different all three times. Auxiliary probes on five endpoints point the same way: short, structured outputs repeat 38–79% of the time, long reasoning outputs 0–36%. The short and long probes differ in form as well as length, so they do not isolate length’s role; they do show that a short check can pass on an endpoint that is not deterministic. Determinism has to be probed with requests that resemble the real workload.
- **A mid-run endpoint shift, and the signal that exposed it.** Between repetition 3 and 4, with every flag unchanged, ten tasks got worse and none got better, and median completion tokens per action fell 20.1% in one step. Several checks on the same trajectories point the same way, and a per-round test fired on the token signal at the boundary round while a score test at the same granularity did not.
- **A noise floor for this setup, and what it implies.** Across the later rounds, two single runs need to differ by 8.07 points before the gap exceeds what repetition alone produces. Resampling whole rounds puts the same quantity on a 95% interval of 5.18–8.53 points, and a shared round-level effect, if one exists, would only raise it. The figure is a property of a model fitted to one agent, one endpoint and one reweighted task sample, not a constant of the benchmark.
- **Findings that need no compute.** Nine checks on the public repository at a pinned commit, each with where to look, that bear on whether two entries on one leaderboard are comparable. Seven require reading code or repository history; two are computed from public trajectories; the last is an inference rather than a read fact. A reader can check each one without trusting any measurement of ours.

A benchmark that reports only a scalar score is a measuring device with no self-check: it cannot distinguish a system that improved from an instrument that moved, and neither can the metadata recorded beside it. Treating the harness as a verification system — recording the signals a failure would show up in, not only the number — is what made Section 4 visible at all, at a cost of one integer per request.

2 Setup: what was pinned

Table 1 is the configuration. Everything in it is a flag we could set. This paper is about what remained free anyway.

Task selection, and why it is skewed on purpose. The 29 tasks are drawn from the benchmark’s own published per-model trajectories: 10 that the published runs disagree on, 9 that all of them pass, and 10 that all of them fail.¹ The disagreed-on group is where within-configuration instability is

¹The groups were assigned from the 15 of the 18 published bundles whose result files our parser read; the other three record results in formats it did not recognise. Read in full, five of the 29 tasks would carry a different group

most likely to appear; the other two are controls, and a setup that produced flips everywhere would be measuring something other than the benchmark. Because the set is deliberately unrepresentative, its absolute score level is not comparable to a leaderboard number. Only widths are, and only after the reweighting in Section 5.

Records kept per run. The harness alone records none of the fields above; what our orchestrator adds, and how every figure traces back to the raw records, is catalogued in Appendix E. The raw records, the endpoint probes behind Section 3 and evidence strands 4 and 5 of Section 4, the analysis code with its derived data files, and the figure scripts are released together at <https://github.com/Universeyi/instrument-moves>.

3 temperature=0 is not determinism

Holding the model still is the first thing a careful experimenter does, and the obvious lever is greedy decoding with a fixed seed. It did not work on either endpoint we tested, and our own first attempt to check that it worked gave the wrong answer.

We sent the same prompt three times to each of two endpoints serving the same model, at temperature=0 with seed=42. Both endpoints advertise seed support in the router’s capability listing. All three completions differed on both endpoints, with lengths varying roughly twofold (1571 / 2812 / 1427 tokens on one, 2963 / 1448 / 1494 on the other).

The trap. An earlier version of this check used the prompt “*Reply with exactly: OK*”. On deepinfra/fp4 it returned byte-identical output twice, and we recorded that the endpoint was deterministic. It is not: three longer requests to the same endpoint all came back different. Auxiliary probes at five endpoints that carry both kinds of request repeat on short, structured outputs 38–79% of the time and on long reasoning outputs 0–36%. The two kinds of request differ in form as well as length, so they do not isolate the role of length, and short outputs are far from always identical. **Determinism has to be probed with requests that resemble the real workload.** We report this because the cheap check is the one a practitioner will reach for, and it failed in the reassuring direction.

Relation to known results. That low-precision inference is non-reproducible across serving configurations is established: Yuan et al. [2025] trace it to the non-associativity of floating-point arithmetic and show that batch size, GPU count and GPU version alone move accuracy by up to nine points under bfloat16 greedy decoding. Our contribution here is not the phenomenon. It is that the phenomenon is invisible to the standard sanity check, and that a hosted endpoint gives a practitioner no way to hold those variables fixed — batch composition is a property of other customers’ traffic.

Consequence for the rest of this paper. Because sampling cannot be held still, we do not claim any decomposition of variance into “environment” and “model” parts. What Section 5 reports is the total variation that survives every control available to a careful experimenter, which is the quantity a reader comparing two leaderboard entries actually faces.

4 The endpoint moved

Ten and a half hours into a thirty-three hour run, between repetition round 3 (which ended 2026-08-26 06:47 UTC) and round 4 (which began 06:49 UTC), the results changed character. Nothing changed on our side: same image digest, same task-set commit, same model identifier, same pinned provider and quantization, same temperature and seed, same code.

4.1 Five checks on the same trajectories, corroborated outside them

1. Every task that moved, moved the same way. Ten tasks changed pass rate across the boundary; all ten got worse, none improved, nineteen were unchanged. Figure 2 shows all ten arrows

label: four *all-pass* tasks (AcceptMeetingTask, CheckEventTimeTask, GoogleMapsAlibabaSouthNeighborTask, TakeSelfieTask) each have one or two failing runs among the extra three, and one *all-fail* task (SMSManagement) has one passing run. The selection and every measurement are unchanged; only the label describing the published runs differs.




Figure 1: Suite score and reasoning volume per repetition round, on a shared axis. Both step at the same place, between rounds 3 and 4 (red line), and both hold at the new level for the remaining seven rounds. The provider outage (shaded, round 6) is a separate event at a different round; the step is not the outage. Round 6 scored 23 of 29 tasks and is annotated accordingly. Two stacked panels rather than one plot with two vertical axes, so that no choice of scale can be doing the work.




Figure 2: Per-task pass rate before and after the boundary, one arrow per task that moved. Ten tasks moved and all ten moved the same way: 10 worse, 0 better, 19 unchanged (exact sign test $p = 0.002$, read as an indication of how one-sided the split is; tasks within a round are not strictly independent).

pointing the same way, which repetition noise alone would rarely produce: an exact sign test on the ten gives $p = 0.002$. We read that as an indication of how one-sided the split is rather than as a hypothesis test at face value, since tasks within a round share a container, the backend services and an execution order, and so are not strictly independent.

2. Reasoning volume stepped, and did not drift. Median completion tokens per action — one step is one model request, and the output is almost entirely reasoning — were 132.5, 132.4, 133.1 in rounds 1–3, within 0.55% of each other, and 102.2, 98.8, 108.7, 105.9, 103.8, 108.4, 108.0 in rounds 4–10, spanning 98.8–108.7, about $\pm 4.8\%$ around the regime midpoint. Between the two regimes the median fell 20.1%.

3. Faster and worse at the same time. Comparing rounds 1–3 with rounds 4–10 throughout (round 6 in, its six zero-step no-verdict runs out): mean time per step fell from 16.0 s to 13.7 s while quality fell. Degrading hardware or a congested network makes an endpoint slower, not quicker. Generating less reasoning per call explains both directions at once, and it explains the third-order

effects too: mean steps per run rose from 28.7 to 30.6 and the share of runs stopped by the 50-step budget rose from 31% to 45%.

Time of day does not explain it. Co-tenant traffic changes batch composition — the mechanism Yuan et al. [2025] identify — so a diurnal effect is a live alternative. Our own schedule rules it out: rounds 9–10 ran 22:21–04:51 UTC, inside the same overnight window as rounds 1–3 (20:10–06:47 UTC), and stayed at the post-shift level. A daily cycle cannot produce a step that survives a return to the original hours.

4 and 5. Corroborated outside the harness, with a control. The determinism probe of Section 3, unchanged, sent to the pinned endpoint on two consecutive days: on the first day the three completions ranged from 1427 to 2812 tokens (sample SD 761), and on the second day they came back nearly identical (1344, 1292, 1292; sample SD 30). The same probe against a second provider serving the same model gave 2963, 1448, 1494 on the first day and 1938, 1672, 2427 on the second, varying on both days. Each cell is three requests, so we describe rather than test, and we report sample SD throughout. The probe locates the change to a day; the round boundary comes from the harness’s internal timing. The harness is not involved in either measurement, and the change in behaviour was specific to the endpoint we had pinned rather than to the model, the prompt or our tooling.

What we observed, and what we infer. We observed an endpoint’s behaviour, not its configuration. A serving-side change is the only explanation we know of that accounts for all of it at once, but we cannot rule out explanations we have not thought of, and nothing below depends on the cause.

4.2 Why this is the finding

Every one of these controls exists and we set it. There is no flag that pins how a provider serves the model, because serving behaviour is not part of the API contract; it is an operational property of somebody else’s fleet. And there is no field in any leaderboard entry, on this board or the others we looked at, that records which configuration produced the number.

This is not a property of the benchmark we happened to use. Any evaluation built on hosted inference is exposed to this gap between what the submitter controls and what determines the number, and the more an agent’s behaviour depends on long reasoning chains, the more exposed it is. We observed one such change in ten repetitions; we have no measurement of how often they happen, and nothing here should be read as a claim that they are routine.

A cheap side signal, offered with its limits. The quantity that made this visible is completion tokens per action. The harness records `total_tokens`, which is 98.5% prompt in this workload and moved by 8% across the change: not enough to notice. Completion tokens per action moved 20.1% in one step, and a per-round check of the kind our sentinel uses fires on it at the boundary round while a score test at the same granularity does not. We call this a side signal, not a detector: it has worked once, after the fact. A sentinel probe logging fixed requests to 14 endpoints over three weeks raised five length alarms in 328 endpoint-days, none on two consecutive days; day-shuffled calibration puts the false-alarm rate at 0.06%. The sentinel measures the output length of fixed requests, a close relative of tokens per action rather than the same quantity, and a quiet period calibrates false alarms, not sensitivity. Recording the signal cost nothing and we had to recover it from raw trajectories after the fact. We recommend logging it by default (Section 7).

5 The resulting noise floor

Seven repetitions fall after the shift in behaviour. They scored 31.0, 31.0, 39.1, 34.5, 37.9, 37.9 and 34.5%. *Every headline figure below is computed on the six rounds that exclude round 6, for the reason given after Table 2; seven-round values are given alongside so the difference is always visible.* Seven of 29 tasks did not return the same verdict every time (eight over all seven rounds), and the controls behave as they should: 0 of 10 flips among tasks all published runs fail, 1 of 9 among tasks all published runs pass, and 6 of 10 among the disagreed-on group.

Table 2: Leave-one-round-out. Every figure recomputed from scratch on what remains. The range responds only to removing the maximum; the same-system difference threshold is stable at 7.91–8.96 points and the comparison against published entries never changes.

round dropped	4	5	6	7	8	9	10
best — worst (pp)	8.10	8.10	6.90	8.10	8.10	8.10	8.10
same-system threshold (pp)	8.96	8.47	8.07	8.38	7.91	8.69	8.96
adjacent gaps within noise	16/17	16/17	16/17	16/17	16/17	16/17	16/17

The headline is not the range. A best-minus-worst spread is an order statistic: it cannot grow when a repetition is removed, it grows on average with the number of repetitions, and a seven-round spread and a six-round spread are therefore not comparable quantities. Our own sensitivity check forced this correction. Table 2 recomputes everything with each round dropped in turn: the spread is unmoved by dropping any round except the one holding the maximum, while the same-system difference threshold varies across all seven variants within a narrow band. Throughout, the *same-system difference threshold at 95%* is $1.96\sqrt{2}\sigma$, with σ the exact standard deviation of a sum of independent per-task Bernoulli outcomes at the reweighted 117-task scale; interval widths beside it are empirical 2.5–97.5% spans of the same distribution, hence not exact multiples (Appendix C).

Round 6, and why the conservative number is the headline. Round 6 contains the only infrastructure incident in the run: the provider returned HTTP 401 for about an hour, 68 of the 95 silent restarts in this regime fall in it (110 across all ten rounds), and all six runs that produced no verdict are in it. It is also the highest-scoring round of the regime, on a reduced denominator of 23 tasks. We therefore report the figures *without* it as the headline, with the seven-round values alongside:²

	6 rounds (round 6 removed)	7 rounds
same-system threshold at 95%	8.07 pp	8.64 pp
95% interval at suite scale, width	11.11 pp	12.82 pp
tasks that flipped	7 / 29	8 / 29

Removing the incident moves the result against us, and we report it anyway. The threshold falls from 8.64 to 8.07 points, which is a stricter bar: the number of published pairs separated by less than it drops from 39 to 35. We report the lower figure because it is the one measured on data containing no infrastructure incident, not because it is the more favourable one.

Two things are worth stating precisely rather than glossing. First, the single task that stops flipping, `DownloadSendReceiptTask`, does so at $1/7 \rightarrow 0/6$: it lost its one passing round, not its instability. Only 4 of 29 tasks change pass rate at all, and no task flips *only* when round 6 is removed. Second, imputing the six lost runs at those tasks’ own rates in the other rounds puts round 6 at 40.2% rather than the 39.13% recorded — the outage cost that round height rather than lending it any, because the tasks it happened to cost pass more often than average (44.4% against 34.5%).

Carrying it to the full suite. The 29 tasks are deliberately skewed, so their level means nothing on its own. Reweighting the three groups to their true sizes in the 117-task GUI-only suite (98 disagreed-on, 9 always-passed, 10 always-failed) gives a width that is comparable to a leaderboard number. The ten disagreed-on tasks we measured are not a random sample: before running anything, and using public data alone, we selected deterministically the ten GUI-only tasks rated by at least ten published runs whose public pass rate was closest to 50% (ties broken by run count, then task name). Tasks near 50% are likely among the least stable, so extrapolating from them probably overstates the width; we state this as a direction, not a measurement. Dropping any single task from the ten moves the threshold between 7.61 and 8.56 points (7.61 without `CVEmailTask`). Stated

²The two scales move in *opposite* directions when round 6 is removed: at the 29-task scale σ rises from 5.87 to 5.92 pp, at the reweighted 117-task scale it falls from 3.12 to 2.91 pp. They are computed differently: the 29-task figure widens because a task losing a repetition widens its Beta posterior, while the 117-task figure narrows because two low-rate tasks move further from 0.5, cutting their Bernoulli variance by more than the two tasks moving toward 0.5 add. The same-system difference threshold derives from the second scale. We report both because reporting only the one that suits the argument is precisely the failure this paper is about.

conditionally, because it is: *if the 88 disagreed-on tasks we did not measure behave like the 10 we did*, a single run of the full suite falls within a 95% interval about 11.1 points wide.

What this means for reading a leaderboard. Under the model above, and for this agent on this endpoint, two single runs need to differ by 8.07 points before the gap exceeds what repetition alone produces. The exact value supersedes the 8.15 of the submitted version, which came from simulation noise. Resampling whole rounds gives a 95% interval of 5.18–8.53 for the same quantity, and even at the low end 11 of the 17 adjacent published pairs stay inside the threshold. Requiring 80% power rather than a 95% exclusion of repetition noise raises the figure to 11.54 points, inside which all 17 adjacent pairs fall. Six rounds cannot rule out a shared round-level effect on top of per-task variance: we observe no excess dispersion (total-score SD 3.08 against 3.98 predicted under independence, $p = 0.70$), but an additive effect of up to 5.25 points per round would still fit the data and would raise the threshold to 16.6. Before the shift, the same computation gives $\sigma = 3.40$ and a threshold of 9.41. Of the 17 adjacent pairs among the 18 published GUI-only entries, 16 are separated by less than 8.07, as are 35 of the 153 pairs overall (39 at the seven-round threshold of 8.64). We state this as an indication of scale and not as a re-scoring: the threshold was measured for one model on 29 tasks and is being applied to entries produced by other agents and other models. The appropriate reading of a gap below the threshold is that it *warrants a second run before citation*, not that the two systems are equal.

Are the flips scoring errors rather than variance? It is the obvious alternative explanation, and we checked it: we inspected six tasks against the task definitions in the public repository, one from the top of a ranking of 111 suspicious runs and five whose failure pattern looked wrong, and found no false positives and no false negatives. Four were genuine agent failures; one was a task defect that did not affect our runs (Table 3); one remains unresolved. Appendix G reports the audit and why we had every incentive to find the opposite.

6 What is already true without running anything

Everything above rests on measurements a reader has to take on trust until the data is inspected. This section does not. Table 3 lists nine checks on the public repository at the pinned commit, each checkable in a few minutes, each bearing on whether two entries on one leaderboard are comparable. The first seven require reading code or repository history; the last two are computed from public trajectories, and the final row is an inference rather than a read fact. We re-verified all nine against commit 83e7b8fc while drafting.

None of this requires trusting a measurement of ours. The first seven rows require reading the repository; the last two are computed from public trajectories. One nuance is empirical rather than read from code: the 30 runs above the cap sit at 51–54, and a plausible innocent cause is a custom agent setting its own bound — whichever it is, the board does not say.

We reported what we found upstream. MobileWorld PR #61 is merged (2026-09-22); PRs #59 and #64 are still open; Sierra’s tau2-bench issue #540 is under discussion. We will update the repository README as these move.

7 Recommendations

These are addressed to benchmark maintainers and to anyone reporting agent evaluations, ourselves included. They are cheap, and most of them are fields rather than experiments.

1. **Record the serving endpoint, not just the model identifier.** Provider and quantization at minimum. Two entries carrying the same model id today can be different weights on different hardware, and a reader has no way to tell. This is the single highest-value change here, and it is one column.

³The version submitted for review reported 90 of the runs at the cap as carrying no score at all. Three of the 18 bundles record results in formats our first parser did not read; read in full, every run at the cap has a score, and the 90 split into 83 zeros and 7 passes.

Table 3: Checkable without running the benchmark. The upper group is read from the code or the repository history, with line numbers at commit 83e7b8fc. The lower group is computed from the benchmark’s own published per-model trajectories; its last row is an inference, and an alternative reading is discussed in the text.

Observation	Where	Consequence for comparability
<i>Read from the code or the repository history</i>		
For model ids containing <code>claude</code> , <code>max_tokens</code> is set to 64000; every other model uses the agent default of 2048.	<code>agents/base.py:97-99</code>	Entries on one board run under generation budgets differing by a factor of 31, and no field records it. Room to reason is what Section 4 shows moving the score; we claim the structural difference only, no entry’s counterfactual.
For model ids containing <code>claude</code> , <code>temperature</code> is deleted from the request.	<code>agents/base.py:97-99</code>	Those entries ran at the provider’s default sampling, not at the value the harness otherwise sets.
For model ids containing <code>kimi-k</code> , <code>extra_body</code> is assigned rather than merged, discarding whatever the caller supplied.	<code>agents/base.py:105-106</code>	Provider and quantization pinning supplied by a caller never reaches the API. This is how we found it: requests we believed pinned were served by three different providers.
The cleanup path matches containers against a hardcoded <code>:latest</code> image.	<code>runtime/utils/models.py:28</code>	With a pinned image tag, <code>mw env rm -all</code> reports success while containers keep running.
The emulator is started with <code>-no-snapshot</code> .	<code>docker/start_emulator.sh:5</code>	The documented guarantee of identical initial state via AVD snapshots is not what executes.
Verifier corrections 8ae5064, c33e340, 32ed916 and 702e904 postdate existing entries.	repository history	Entries from different dates were scored by different verifier versions, and no entry records which.
ReadQwen3PaperTask5 scores by exact string match against <code>vie_Latn_khm Khmr</code> ; only whitespace around commas is normalised.	<code>tasks/definitions/native/read_paper_5.py:21, 89-96</code>	A correct answer in plain English, in the other order, or in the notation the source paper prints (<code>vie_Latn</code>) scores 0. Our ten runs of this task failed earlier, with an empty answer, so no result here depends on it.
<i>Computed from the published trajectories</i>		
Of 2562 published (model, task) runs with a step count, 525 (20.5%) sit at exactly 50 steps, against 18, 10, 20 and 2 at 48, 49, 51 and 52. 490 of the 525 (93.3%) scored 0 and 35 passed. ³	published trajectories	One run in five ends at a budget rather than at a task outcome; the record shows a zero, not which event occurred.
All 30 runs above 50 steps belong to a single entry, which has no runs at exactly 50; the other 17 entries have runs at 50 and none above.	published trajectories	<i>Inference:</i> consistent with at least two different round budgets behind the published board, and no entry records which applied.

2. **Record image digest, task-set commit and run date.** Digest rather than tag: a tag moves. The run date is what lets a reader notice that two entries are separated by an unannounced change in endpoint behaviour.
3. **Report the number of runs, and a range when there is more than one.** Three runs with a minimum and maximum costs three times the compute and removes most of the ambiguity. Where a board already stores a run count, surfacing it in the rendered table is free.
4. **State a minimum meaningful gap in the leaderboard documentation.** The same-system difference threshold of Section 5 is one way to compute one; resampling whole rounds puts it between 5.18 and 8.53 points. Differences below that range between single runs are not resolvable. Saying so protects submitters as much as readers: it removes the incentive to claim a lead that a re-run would erase.

5. **Log completion tokens per action.** It is the cheapest available witness to an endpoint changing underneath a run. Aggregate token counts are dominated by the prompt and will not show it.
6. **Distinguish budget exhaustion from task failure in the recorded outcome.** A run stopped at the round cap and a run that did the wrong thing are different events; one in five published runs is the former.
7. **Three small fixes we can offer upstream,** each independent of everything above: merge rather than assign caller-supplied request options; guard the response-content accessor against a None body, which a reasoning model returns when it exhausts its budget mid-thought; and match container cleanup against the configured image rather than a hardcoded tag.

There is a stronger conclusion available, and our data supports it: if a serving configuration is not part of an API contract, then an evaluation built on hosted inference is not reproducible in principle, only in practice and only until the provider changes something. The strict remedy is to self-host and freeze the weights. We are not proposing that everyone do so — for most groups the cost is prohibitive and the hosted endpoint is the thing they actually deploy against — but a field that cites two-point differences should be explicit that it has chosen convenience over reproducibility, rather than assuming it has both.

None of this makes a benchmark deterministic, and none of it should. The purpose is narrower: to let a reader of a leaderboard tell the difference between a system that improved and an instrument that moved.

8 Limitations

One agent, one model, one endpoint — and one event. Instability is very likely model-dependent: a model that is decisive on a task will flip less than one that is borderline on it. The threshold we report should be read as an order of magnitude for this class of benchmark, not as a constant to be applied to other entries as though it were measured on them. Separately, we observed one behavioural shift in ten repetitions and have no measurement of how often such shifts occur, so the cost and benefit of the mitigations we recommend cannot be estimated from our data.

Seven repetitions over 29 tasks. Per-task rates are correspondingly imprecise — a task passing 3 of 6 times has a 95% interval of roughly 19–81%. The suite-level figures are much better determined than the per-task ones, and we lean on the former.

The task set is deliberately skewed toward tasks that published runs disagree on. Its score level is not comparable to a leaderboard number; only the width is, and only after reweighting, and that reweighting assumes the 88 disagreed-on tasks we did not run behave like the 10 we did.

GUI-only. The user-interaction and MCP suites were not run. The former puts a second language model inside the loop and would plausibly be noisier, not less.

The harness absorbs part of what we set out to measure. Its internal retry fired 110 times across all ten rounds (95 of them within the seven rounds of the headline configuration), silently, and those restarts are instability that never reached a score. Our figures are therefore a floor.

Environment version. We ran a container image newer than the one most published entries used, because the older tag carries a defect the maintainers had already fixed and injecting a known bug into a variance measurement would measure the wrong thing. The quantization also differs from the reference entry's. Both are recorded per run.

The behavioural shift is documented, its cause is not. We can show when the endpoint's behaviour changed, that the change was specific to one endpoint, and what it did to the output. We cannot see inside somebody else's fleet, and we do not speculate about the cause.

The side signal is calibrated on one event. Completion tokens per action exposed the shift of Section 4, but only in hindsight and only once. The sentinel run gives false-alarm rates, not sensitivity: in a three-week window with no known changes it can say how often it would alarm by chance, not how often a real change would trip it. And the sentinel measures the output length of fixed requests,

which is related to, but not the same as, tokens per action in a real workload. We offer the signal because it costs a log line, not because we have validated it as a detector.

The raw per-run records, the analysis code and the figure scripts are released, and the figures in this paper regenerate from a single command. If the noise floor of an evaluation instrument is worth knowing — and we think a field that cites two-point differences should want to know it — then it is worth publishing alongside the scores, by whoever is best placed to measure it.

9 Related work

Mobile and GUI agent benchmarks. Interactive Android environments with programmatic reward signals were established by AndroidWorld [Rawles et al., 2025], which supplies tasks across real applications on a live emulator rather than a static test set. MobileWorld [Kong et al., 2026], the benchmark studied here, extends that setting with agent–user interaction and MCP-augmented tasks and publishes a leaderboard together with per-model trajectories. Our use of it is not adversarial: the published trajectories are what made the zero-compute analysis of Section 6 possible at all, and few benchmarks release enough for an outsider to do it.

Nondeterminism in language model inference. Yuan et al. [2025] establish that low-precision inference is not reproducible across serving configurations, tracing it to the non-associativity of floating-point arithmetic and showing that batch size, GPU count and GPU version alone can move accuracy by several points under greedy decoding. That work explains the mechanism behind our Section 3 and, plausibly, behind Section 4. The difference in setting matters: their variables are controllable by whoever owns the hardware, whereas a practitioner calling a hosted endpoint controls none of them and cannot observe them either — batch composition is a function of other customers’ traffic.

Drift in hosted model services. That a commercial model endpoint changes behaviour over time is established. Chen et al. [2024] evaluated the March and June 2023 releases of two hosted models across several task families and found large swings — identification of prime versus composite numbers fell from 84% to 51% accuracy on one of them — and argued for continuous monitoring. The same group had earlier framed the general problem of detecting shifts in machine learning APIs efficiently [Chen et al., 2021].

Our setting is narrower and, we think, harder to defend against. That work compares *labelled versions* of a service *months apart*, and detects the change by periodically re-running a whole evaluation suite. What we record happened inside a single model identifier, at one provider, at one quantization, within *a single half-day*, with no version change announced or announceable — and it was caught not by re-running an evaluation but by a cheap side signal, completion tokens per action, that costs one integer per request to log. The first result tells a practitioner that a service will change. The second tells them it can change *during their experiment*, while every control they know about is still set, and that the standard check will not show it.

Flaky tests. Software engineering has a mature literature on tests that pass and fail without a code change, including how practitioners experience and manage them [Gruber and Fraser, 2022]. The analogy is close and the vocabulary is useful, but the disanalogy is the point of this paper: a flaky test is understood as a defect to be fixed, whereas run-to-run variance in an environment-driven agent benchmark is a property of the measurement that should be quantified and reported rather than eliminated. The software engineering framing also assumes the system under test is the thing that varies. Here the harness, the environment and the serving endpoint all vary, and only one of them is the object of measurement.

Statistical reporting for evaluations. Miller [2024] sets out error-bar methodology for LLM evals, Bowyer et al. [2025] caution that CLT intervals understate uncertainty at task counts like ours, and τ -bench [Yao et al., 2025] reports run-to-run inconsistency as a metric, pass^k . All three hold the thing measured stationary while it is measured. Section 5 is what their machinery yields here; Section 4 is the case none of them models — the instrument moving mid-experiment.

Positioning. These lines of work measure model or agent capability, or the reliability of a test suite. This paper measures the instrument: how far a benchmark’s own output moves when its input is held as fixed as the available controls allow, and what a reader would need recorded in order to notice when it moves for reasons that have nothing to do with the agent.

References

- Sam Bowyer, Laurence Aitchison, and Desi R. Ivanova. Position: Don’t use the CLT in LLM evals with fewer than a few hundred datapoints. In *International Conference on Machine Learning (ICML)*, 2025. Spotlight position paper; arXiv:2503.01747; dblp conf/icml/BowyerAI25.
- Lingjiao Chen, Tracy Cai, Matei Zaharia, and James Zou. Did the model change? Efficiently assessing machine learning API shifts. *arXiv preprint arXiv:2107.14203*, 2021.
- Lingjiao Chen, Matei Zaharia, and James Zou. How is ChatGPT’s behavior changing over time? *Harvard Data Science Review*, 6(2), 2024. arXiv:2307.09009.
- Martin Gruber and Gordon Fraser. A survey on how test flakiness affects developers and what support they need to address it. In *IEEE International Conference on Software Testing, Verification and Validation (ICST)*, 2022. arXiv:2203.00483.
- Quyu Kong, Xu Zhang, Zhenyu Yang, Nolan Gao, Chen Liu, Panrong Tong, Chenglin Cai, Hanzhang Zhou, Jianan Zhang, Liangyu Chen, Zhidan Liu, Steven Hoi, and Yue Wang. MobileWorld: Benchmarking autonomous mobile agents in agent-user interactive and MCP-augmented environments. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026. arXiv:2512.19432; dblp conf/acl/KongZYGLTCZZCLH26.
- Evan Miller. Adding error bars to evals: A statistical approach to language model evaluations. *arXiv preprint arXiv:2411.00640*, 2024.
- Christopher Rawles, Sarah Clinckemaillie, Yifan Chang, Jonathan Waltz, Gabrielle Lau, Marybeth Fair, Alice Li, William Bishop, Wei Li, Folawiyo Campbell-Ajala, Daniel Toyama, Robert Berry, Divya Tyamagundlu, Timothy Lillicrap, and Oriana Riva. AndroidWorld: A dynamic benchmarking environment for autonomous agents. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2405.14573; dblp conf/iclr/RawlesCCWLFLBLC25.
- Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2406.12045; dblp conf/iclr/YaoSRN25.
- Jiayi Yuan, Hao Li, Xinheng Ding, Wenya Xie, Yu-Jhe Li, Wentian Zhao, Kun Wan, Jing Shi, Xia Hu, and Zirui Liu. Understanding and mitigating numerical sources of nondeterminism in LLM inference. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. Oral presentation; arXiv:2506.09501.

A Per-task results

All 29 tasks over the six repetitions of the headline basis, ordered by flip rate. **The group column describes the published runs’ behaviour, not ours:** it records whether the benchmark’s own published trajectories all passed a task, all failed it, or disagreed on it (Section 2), and it is how the task set was stratified. A task can therefore be listed as *all-pass* and still have failed in every one of our repetitions.⁴ Flip rate is the share of trial pairs that disagree, $k(n - k) / \binom{n}{2}$; its maximum is below 1 and depends on n , reaching 0.600 at $n = 6$. Wilson intervals are wide at this sample size, which is why the suite-level figures rather than the per-task ones carry the argument in Section 5.

⁴Four tasks here score 0/6 despite having passed in two of the three rounds that preceded the shift in behaviour of Section 4: BidFileRenameTask, CheckEventTimeTask, DownloadSendReceiptTask and MattermostCreateChannelTask (see Figure 2). The headline basis lies entirely after that change, so these rows are a consequence of it rather than an anomaly. The remaining zero-scoring tasks are ones the published runs also fail.

task	group	passed	pass rate	Wilson 95%	flip rate	steps (median)
CVEmailTask	disagreed	3/6	50%	19–81%	0.600	11
ScheduleLunchViaSmsTask	disagreed	4/6	67%	30–90%	0.533	25
MattermostBudgetApprovalPipelineTask	disagreed	2/6	33%	10–70%	0.533	50
MattermostEmailTask	disagreed	2/6	33%	10–70%	0.533	21
TakeSelfieTask	all-pass	5/6	83%	44–97%	0.333	12
MastodonImportMutedUsersTask	disagreed	1/6	17%	3–56%	0.333	50
SendFormsTask	disagreed	1/6	17%	3–56%	0.333	50
AcceptMeetingTask	all-pass	6/6	100%	61–100%	0.000	7
ChangeWallpaperTask	all-pass	6/6	100%	61–100%	0.000	9
CloseFlightModeTask	all-pass	6/6	100%	61–100%	0.000	7
GoogleMapsAlibabaSouthNeighborTask	all-pass	6/6	100%	61–100%	0.000	9
MastodonChangeHeaderTask	all-pass	6/6	100%	61–100%	0.000	7
MastodonNewPostTask	all-pass	6/6	100%	61–100%	0.000	6
OpenFlightModeTask	all-pass	6/6	100%	61–100%	0.000	6
BidFileRenameTask	disagreed	0/6	0%	0–39%	0.000	50
CheckEventTimeTask	all-pass	0/6	0%	0–39%	0.000	20
CheckInterviewTimesTask	all-fail	0/6	0%	0–39%	0.000	50
DownloadSendReceiptTask	disagreed	0/6	0%	0–39%	0.000	34
GraduationMassEmailTask	disagreed	0/6	0%	0–39%	0.000	13
LocalFileManagementTask2	all-fail	0/6	0%	0–39%	0.000	50
MastodonCalendarMultiMemosTask	all-fail	0/6	0%	0–39%	0.000	50
MastodonGetServerInfoTask	all-fail	0/6	0%	0–39%	0.000	50
MastodonServerInfoReportTask	all-fail	0/6	0%	0–39%	0.000	50
MastodonShareLocationTask	all-fail	0/6	0%	0–39%	0.000	36
MattermostCreateChannelTask	disagreed	0/6	0%	0–39%	0.000	24
MattermostProjectStatusReportTask	all-fail	0/6	0%	0–39%	0.000	50
MattermostResourceConflictResolutionTask	all-fail	0/6	0%	0–39%	0.000	50
ReadQwen3PaperTask5	all-fail	0/6	0%	0–39%	0.000	50
SMSManagement	all-fail	0/6	0%	0–39%	0.000	50

B Tasks that never flipped

The 22 tasks that returned the same verdict in every repetition of the headline basis. As in Appendix A, *group* is the published runs’ behaviour rather than ours, which is why a row can read *all-pass* and *always failed* at once. Their existence is the control: a setup that flipped everywhere would be measuring something other than the benchmark. Note that all 10 tasks the published runs uniformly fail are here, and 8 of the 9 they uniformly pass.

task	group	passed	outcome
CheckInterviewTimesTask	all-fail	0/6	always failed
LocalFileManagementTask2	all-fail	0/6	always failed
MastodonCalendarMultiMemosTask	all-fail	0/6	always failed
MastodonGetServerInfoTask	all-fail	0/6	always failed
MastodonServerInfoReportTask	all-fail	0/6	always failed
MastodonShareLocationTask	all-fail	0/6	always failed
MattermostProjectStatusReportTask	all-fail	0/6	always failed
MattermostResourceConflictResolutionTask	all-fail	0/6	always failed
ReadQwen3PaperTask5	all-fail	0/6	always failed
SMSManagement	all-fail	0/6	always failed
AcceptMeetingTask	all-pass	6/6	always passed
ChangeWallpaperTask	all-pass	6/6	always passed
CheckEventTimeTask	all-pass	0/6	always failed
CloseFlightModeTask	all-pass	6/6	always passed
GoogleMapsAlibabaSouthNeighborTask	all-pass	6/6	always passed
MastodonChangeHeaderTask	all-pass	6/6	always passed
MastodonNewPostTask	all-pass	6/6	always passed
OpenFlightModeTask	all-pass	6/6	always passed
BidFileRenameTask	disagreed	0/6	always failed
DownloadSendReceiptTask	disagreed	0/6	always failed
GraduationMassEmailTask	disagreed	0/6	always failed
MattermostCreateChannelTask	disagreed	0/6	always failed

C Three intervals that are easy to confuse

We compute and report three different quantities, and conflating them is the most likely way to misread this paper.

- **Re-run distribution.** Where the total would land if the same task set were run again. Each task’s rate is drawn from a Jeffreys posterior before the Bernoulli draw, so that the finite

number of repetitions is carried through rather than assumed away. This is the headline at the 29-task suite scale (the σ of the footnote in Section 5).

- **Plug-in distribution.** The same, with each measured rate treated as exact. The reweighted 117-task width and the same-system difference threshold — $1.96\sqrt{2}\sigma$ of this distribution — derive from it. Excluding parameter uncertainty makes both lower bounds on the re-run width, which is the conservative direction for reading a gap.
- **Estimate bootstrap.** How precisely *this experiment* determined its own number, by resampling repetitions within each task. It answers a question about our precision, not about a future run.

D Silent restarts

Across the 290 runs of the arm the harness’s in-task retry fired 110 times. Every one is instability that was absorbed before it could reach a score, which is why the variance reported in Section 5 is a floor. The concentration is the diagnostic: a hard task is hard in every round, whereas an incident happens in one.

round	1	2	3	4	5	6	7	8	9	10
silent restarts	6	5	4	3	3	68	3	7	7	4

By task, counting only tasks that were restarted at least once:

task	silent restarts
MastodonImportMutedUsersTask	13
MastodonServerInfoReportTask	12
MastodonShareLocationTask	12
AcceptMeetingTask	9
CVEmailTask	9
ChangeWallpaperTask	9
MattermostCreateChannelTask	9
MattermostEmailTask	9
ScheduleLunchViaSmsTask	9
SendFormsTask	9
MattermostBudgetApprovalPipelineTask	5
MastodonGetServerInfoTask	4
DownloadSendReceiptTask	1

E Method and provenance

The harness records a trajectory, a screenshot per step and a free-text verdict, but no timestamps, model parameters, image digest or task-set commit. Our orchestrator adds those, plus the count of directories the harness leaves behind when it silently restarts a run, and writes one `fsync'd` line per run. The raw record (`runs.jsonl`, 314 lines covering this arm and the preparatory smoke and calibration arms) carries a `sha256` that is stamped into the analysis file and onto every report figure; the paper’s figures regenerate from that analysis file by one seeded command, so any figure can be traced to the records that produced it.

Every field below is recorded on each of the 290 rows and checked for consistency across them; a row disagreeing with the others would be reported as a violation rather than averaged over.

model	moonshotai/kimi-k2.5
serving provider / quantization	atlas-cloud / int4
sampling	temperature 0.0, seed 42
agent	agents/pinned_e2e_agent.py
image	ghcr.io/tongyi-mai/mobile_world:v1.4
image digest	sha256:00e24a8995892af9a0e1b68bc14ce455 fae288d0f1ca8cf57e84a58e40e9f19e
task-set commit	83e7b8fc75ebb6c4a098254a42999db5f4172666
max rounds / step wait	50 / 3.0s
harness auto-retry	10
container recycling	trial
host kernel / cores	6.17.0-1022-gcp / 8
pass threshold	score > 0.99
resampling	10000 draws, seed 20260824
rows analysed (headline basis)	174 (rounds 4–10 less round 6; 203 over all seven, 290 over all ten)
duplicate trial keys	0 (none expected)
provenance consistent across rows	yes

Round boundaries in wall-clock time, so that the change of Section 4 and the outage of Section 5 can be located independently of the prose:

round	start (UTC)	end (UTC)	hours into the run
1	08-25 20:10	08-25 23:17	0.0
2	08-25 23:19	08-26 03:11	3.2
3	08-26 03:13	08-26 06:47	7.1
4	08-26 06:49	08-26 09:53	10.7 ← change
5	08-26 09:55	08-26 13:05	13.8
6	08-26 13:07	08-26 16:25	17.0 ← outage
7	08-26 16:27	08-26 19:24	20.3
8	08-26 19:26	08-26 22:18	23.3
9	08-26 22:21	08-27 01:30	26.2
10	08-27 01:33	08-27 04:51	29.4

The evaluation, the analysis and the figures are three commands:

```
python3 scripts/run_matrix.py --arm A_temp0 --trials 10
--tasks data/pilot_all.txt --agent-type <pinned agent>
--temperature 0 --seed 42
python3 scripts/analyze.py --arm A_temp0 --trials 4:10
--exclude-trials 6 --leave-one-out --regime-split 4
python3 scripts/plots_paper.py
```

Resampling is seeded, so the same raw records always produce the same analysis file, and the analysis file carries the sha256 of the records it was computed from. A separate script re-checks the paper’s headline numeric claims against those files and fails if any disagrees, or if the raw-record hash appears in the text.

One class of number does not come from that pipeline: the endpoint probes of Section 3 and evidence 4–5 of Section 4 were produced by a standalone diagnostic against the live router and transcribed into a data file, which records that provenance explicitly. We flag the difference rather than let it pass as equivalent.

F Tasks that randomise their own initial state

Seven tasks in the benchmark call the standard library’s random functions in their own setup, and no seeding call appears anywhere in the codebase. For these, successive runs are not literally the same problem, so a flip carries a different meaning: the content changed, not only the outcome.

One of the seven, CVEmailTask, is in our task set. **It is included in the 29-task aggregate**, and it flips (3/6 on the headline basis). We flag it rather than drop it: excluding a task after seeing that it flips would be selecting on the outcome, and one task cannot move a 29-task total by more than

3.4 points in any case. The analysis records it as a separate subgroup so that a reader who prefers to discount it can, and Appendix A identifies it by name.

G Checking the alternative explanation: verifier errors

If the flips in Section 5 were scoring mistakes rather than variance, the paper’s central quantity would be measuring the wrong thing. We therefore looked, and we had every incentive to find something: a demonstrable verifier defect would have been a stronger and more quotable result than the one we have, and it would have fitted this paper’s theme better.

We ranked runs by how likely they were to be scoring errors — minority outcomes inside unstable tasks first, then passes completed in fewer steps than any published run of the same task needed, then runs stopped at the round budget. The top of that ranking led to one task; we added five whose failure pattern looked wrong: an identical or implausible failure reason on every trial, or a flip among tasks every published run passes. We inspected these six tasks against the task definitions in the public repository.

Among the six we found no false positives and no false negatives. Four were genuine agent failures (MattermostEmailTask, MastodonGetServerInfoTask, MastodonShareLocationTask, Check-EventTimeTask). One was a task defect: ReadQwen3PaperTask5 requires an exact match to `vie_Latn, khm Khmr`, and the English answer, a reversed order, and the paper’s own spelling `vie_Latn` all score zero. Our ten runs all failed earlier, with empty answers, so the defect did not affect our measurements; it is listed in Table 3. One remains unresolved: GraduationMassEmailTask, where published data disagree, a published Kimi-K2.5 entry passes, and our ten runs all failed with the same “No email found” error, which we did not pursue further.

We report the result for three reasons: it is the obvious alternative explanation and a reader deserves it answered before asking; it leaves that explanation unsupported where it was most likely to show, which is what Section 5 needs — six tasks is a triage depth, not an exhaustive audit, and the ranking is released so a reader can extend the search from the top; and reporting a search that failed is the discipline this paper argues for.